

A Unified Homotopy Framework for Nonlinear System Identification and Simulation: Constructive Dimensioning of Bilinear Neural Architectures

Rodolfo H. Rodrigo^{a,*}, H. Daniel Patiño^a

^a*Instituto de Automática (INAUT), UNSJ-CONICET, Av. Libertador San Martín
(Oeste) 1109, San Juan, J5400ARL, San Juan, Argentina*

Abstract

We develop a unified algorithmic framework for nonlinear system identification and forward simulation of ordinary differential equations, built on Liao’s zeroth-order homotopy deformation equation. Both computational tasks—fitting a parametric function approximator from input–output data, and forward-integrating the resulting nonlinear ODE—are reduced to closed-form Newton (z_1) and Halley (z_2) corrections derived from the same deformation equation. The framework exploits the bilinear parameter structure of common function representations (radial basis functions, sigmoid multilayer perceptrons): a linear last-layer that admits an exact one-step least-squares solution, and nonlinear internal parameters that admit curvature-aware corrections without a scalar learning rate.

The principal theoretical contribution is a constructive dimensioning chain $\varepsilon_{\text{sim}} \rightarrow \varepsilon_{\text{id}} \rightarrow M_{\text{min}} \rightarrow N_{\text{min}} \rightarrow T_{\text{max}}$ that maps a target forward-simulation accuracy to the minimum number of basis functions, the minimum number of training samples, and the maximum admissible integration time step. The chain is constructive in the sense that each link is computable from approximation-theoretic bounds and system quantities. An operator-theoretic interpretation emerges: the choice of the auxiliary linear operator $\mathcal{L}$ encodes prior knowledge that reduces the residual nonlinearity the approximator must

*Corresponding author

Email addresses: rrodrigo@inaut.unsj.edu.ar (Rodolfo H. Rodrigo),
dpatino@inaut.unsj.edu.ar (H. Daniel Patiño)

learn, and therefore reduces the required approximator dimension.

We validate the framework on benchmark problems: a one-dimensional discharge-flow system (orifice-tank model), a two-dimensional Lotka-Volterra predator-prey system, and robustness experiments under measurement noise. Gaussian RBF and sigmoid MLP approximators identified from sparse data (20–40 points) attain identification residuals from $3 \cdot 10^{-2}$ (RBF) down to machine precision ($1.7 \cdot 10^{-8}$, MLP) in a single outer iteration and yield forward-simulation errors of at most 2.3% over 200-step horizons on the 1D benchmark. The scope of the empirical validation is deliberately narrow—low-dimensional ODEs, small-scale approximators—and the role of the present work is to establish the algorithmic framework, the dimensioning theory, and the proof-of-concept computational regime; broader empirical assessment and extension to high-dimensional or stiff systems are identified as follow-up work.

Keywords: Homotopy analysis method, nonlinear system identification, ordinary differential equations, radial basis functions, Newton–Halley corrections, constructive dimensioning, bilinear approximation, function approximation theory, computational methods

1. Introduction

Consider a nonlinear dynamical system observed through input–output data,

$$\dot{y}(t) + f(y(t)) = u(t), \quad y(0) = y_0, \quad (1)$$

where f is an unknown nonlinearity and $u(t)$ is a measured excitation. Two classical computational problems arise: *identification*—reconstruct f from data $\{(t_i, y_i, u_i)\}_{i=1}^N$ using a parametric function approximator—and *forward simulation*—given an identified $\hat{f}$, integrate (1) forward for new inputs and initial conditions. The two are usually attacked with separate computational machinery. Identification employs gradient-based optimization of parametric approximators (radial basis functions, multilayer perceptrons [20, 35, 43, 34]) or, more recently, sparse symbolic regression [9, 13]. Simulation employs explicit and implicit time-stepping schemes [19, 11]. Each approach has its own convergence theory, tuning parameters, and failure modes.

This paper isolates an algebraic structure that is common to both problems and exploits it for joint, hyperparameter-light treatment. The starting point is the zeroth-order deformation equation of the Homotopy Analysis

17 Method [26, 28, 27],

$$(1 - q) \mathcal{L}[\phi(q) - \phi_0] = q c_0 \mathcal{N}[\phi(q)], \quad q \in [0, 1], \quad (2)$$

18 in which $\mathcal{L}$ is a chosen linear operator, ϕ_0 is an initial approximation, and c_0
 19 is a convergence-control parameter. Equation (2) is agnostic to the nature of
 20 ϕ and $\mathcal{N}$: ϕ may be a parameter vector and $\mathcal{N}$ an approximation residual,
 21 or ϕ may be a state trajectory and $\mathcal{N}$ a discretised ODE residual. Truncat-
 22 ing the homotopy series and matching powers of q yields three closed-form
 23 corrections z_1, z_2, z_3 whose algebraic shape depends only on the local Taylor
 24 coefficients of $\mathcal{N}$ at ϕ_0 , not on the dimension or interpretation of ϕ .

25 *Contributions.* The main contributions are the following.

- 26 (C1) *Universal deformation algebra.* We show that the corrections z_1 and z_2 ,
 27 when read off the deformation equation (2), recover Newton’s method
 28 (with the choice $\lambda = \mathcal{N}'_0$, $\sigma = -1$, $z_2 = 0$) and Halley’s method as
 29 instances, and admit a parametric extension to arbitrary λ . This places
 30 HAM, Newton, and Halley in a common scalar parametrisation (γ, σ)
 31 that is exploited uniformly across identification, simulation, and the
 32 linear/nonlinear sub-problems within identification.
- 33 (C2) *Bilinear factorisation of identification.* For both RBF and MLP ap-
 34 proximators, the identification residual is linear in the last-layer weights
 35 and nonlinear in the internal parameters. The framework solves the lin-
 36 ear sub-problem exactly in one step (z_1 with quadratic form, $z_2 = z_3 =$
 37 0 structurally) and applies homotopy corrections only to the nonlin-
 38 ear sub-problem. This is a Variable-Projection-style separation [17, 44]
 39 but with curvature-aware corrections instead of gradient descent on the
 40 projected residual.
- 41 (C3) *Derivative-free integral simulation.* For the simulation step we recast
 42 (1) as a Volterra integral equation with exponential kernel $K(t, s) =$
 43 $e^{-k(t-s)}$, where k is the linearisation of the identified nonlinearity at
 44 the operating point. The discrete residual then has all coefficients of
 45 order $O(1)$ and contains no $1/T$ factor; sensor noise is attenuated rather
 46 than amplified. The homotopy corrections on this residual give per-step
 47 convergence in 1–2 evaluations.

48 (C4) *Constructive dimensioning.* We give explicit conditions linking target
 49 simulation error ε_{sim} to identification error ε_{id} (Gronwall propagation),
 50 to the minimum number of basis functions M_{min} (RBF approximation
 51 rate à la Buhmann 2003; Barron norm 1993 for sigmoid MLPs), to the
 52 minimum number of training samples N_{min} (Jacobian conditioning),
 53 and to the maximum admissible step T_{max} (Newton convergence on the
 54 integral residual). The chain is constructive in the sense that each link
 55 is computable from system quantities and the others are determined.

56 (C5) *Inductive bias as operator choice.* The dimensioning chain shows that
 57 M_{min} depends not on $\|f''\|_{\infty}$ globally but on $\|\tilde{f}''\|_{\infty}$, where $\tilde{f} = f -$
 58 $\mathcal{L}^{-1}[f]$ is the residual nonlinearity left after the linear operator absorbs
 59 the part of f that is already known. Better prior knowledge produces
 60 a smaller $\tilde{f}$ and reduces the network size needed for a given accu-
 61 racy. This formalises a quantitative meaning of “inductive bias” that
 62 is operator-side rather than loss-side, in contrast to Physics-Informed
 63 Neural Networks (PINNs) [38] where physics enters the loss as a soft
 64 penalty.

65 *Scope of the empirical study.* The numerical results in Section 9 cover a one-
 66 dimensional first-order ODE (orifice-tank discharge) with two architectures
 67 (5-centre RBF, 8-unit sigmoid MLP) and 20 training points. The role of
 68 these experiments is to verify, in the simplest non-trivial setting, that the
 69 closed-form corrections converge as predicted, that the linear/nonlinear sep-
 70 aration behaves as analysed, and that the dimensioning chain is consistent
 71 with the observed network sizes. Generalisation to multi-dimensional, stiff,
 72 or chaotic dynamics, and to convolutional, recurrent, or attention-based ar-
 73 chitectures, is consistent with the bilinear structure but requires follow-up
 74 work; we discuss this explicitly in Section 10.

75 *Outline.* Section 2 positions the contribution against gradient-based training,
 76 Newton-type methods for neural networks, PINNs and Neural ODEs, and
 77 symbolic regression. Section 3 states the deformation equation and derives
 78 the universal scalar corrections. Section 4 develops the identification level
 79 for RBF and MLP architectures. Section 5 develops the integral simulation
 80 level. Section 6 states the unification result in compact form. Section 7 es-
 81 tablishes the constructive dimensioning chain. Section 8 interprets the chain
 82 as a formalisation of inductive bias. Section 9 reports the numerical demon-
 83 strations. Section 10 delineates the empirical scope and identifies follow-up

84 work. Section 11 clarifies the relationship to companion manuscripts by the
85 authors. Section 13 concludes.

86 2. Related work

87 *Gradient-based neural network training.* Backpropagation [42] computes the
88 gradient of a scalar loss with respect to all parameters in a single backward
89 pass; first-order methods such as stochastic gradient descent and Adam [1]
90 then update the parameters via a learning-rate-scaled step in the negative-
91 gradient direction. The combination is the workhorse of modern deep learn-
92 ing, but for the small networks and physically-interpretable problems we con-
93 sider here it has well-known disadvantages: a learning rate must be tuned,
94 convergence is at best linear in the loss, and the quality of the final iterate
95 depends on initialisation and stopping.

96 *Newton-type and second-order methods for neural networks.* Second-order
97 methods exploit curvature information through the Hessian or an approx-
98 imation thereof. Becker and LeCun [5] proposed a diagonal Hessian update
99 for MLPs. Martens [31] introduced Hessian-free optimisation, which builds
100 Newton steps via conjugate gradient on Hessian-vector products. The K-FAC
101 family [32] approximates the Hessian by a Kronecker-factored block-diagonal
102 structure. The framework of the present paper is closer in spirit to Sjölund
103 et al. [44], who exploit the linear/nonlinear separation of single-hidden-layer
104 networks via Variable Projection. The principal differences are: (i) we de-
105 rive the corrections z_1, z_2 from the deformation equation rather than from a
106 Levenberg–Marquardt or Gauss–Newton starting point, which gives a uni-
107 form parametrisation across identification and simulation; (ii) we apply the
108 same corrections to ODE simulation, not only to identification; and (iii) we
109 add an explicit dimensioning analysis that links the corrections’ convergence
110 conditions to network size and data sufficiency.

111 *Physics-Informed Neural Networks.* PINNs [38, 23] embed the governing dif-
112 ferential equation into a soft penalty in the training loss and minimise the
113 combined data-plus-physics loss with backpropagation. The architecture is
114 left untouched. Two consequences follow: (i) the network size must be cho-
115 sen heuristically because the physics term does not change the function class,
116 only the search direction; (ii) the loss landscape becomes ill-conditioned when
117 the physics weight is large or the system is stiff [47, 25]. The framework of

the present paper places physics in the auxiliary linear operator $\mathcal{L}$ and consequently in the residual nonlinearity $\tilde{f} = f - \mathcal{L}^{-1}[f]$ that the network must approximate. The effect on architecture is constructive: better prior knowledge shrinks $\|\tilde{f}''\|_\infty$ and therefore the required M .

Neural Ordinary Differential Equations. Chen et al. [15] parametrise the right-hand side of an ODE by a neural network and train it via the adjoint method. This is conceptually orthogonal to our work: their primary problem is to back-propagate through the ODE solver efficiently for general loss functions on the trajectory; our primary problem is to identify f from sparse data and then simulate accurately. Rackauckas et al. [37] unify these directions under the Universal Differential Equations programme, where physics-known terms are combined with neural-network-learned residual terms; that combination is precisely the operator-side decomposition $f = \mathcal{L}^{-1}[\cdot] + \tilde{f}$ that we use here, and our dimensioning analysis applies to it.

Symbolic and sparse identification. SINDy [9] and its extensions [13, 22] reconstruct f from data by sparse regression on a candidate library of nonlinear basis functions, with ℓ_1 -style penalties to enforce parsimony. The output is a symbolic expression rather than a black-box approximator. Our framework is intermediate between SINDy and classical neural identification: the basis is a smooth parametrised family (RBF, MLP) rather than a fixed library, but the linear sub-problem is solved exactly as in SINDy and the nonlinear corrections are closed-form rather than gradient-based.

Reservoir computing and echo-state networks. Lukoševičius and Jaeger [29] train only the linear output layer of a recurrent network, freezing the nonlinear hidden dynamics. The linear/nonlinear split is identical in spirit to the one we use, restricted to the case in which the nonlinear part is fixed by design. Our framework treats both halves: the linear part exactly, the nonlinear part by curvature-aware corrections.

Homotopy methods in numerical analysis and control. The Homotopy Analysis Method [26, 28] was developed for analytic series solutions to nonlinear differential and integral equations. Turkyilmazoglu [45] and Vajravelu and Van Gorder [46] analysed convergence properties and the role of the auxiliary parameter c_0 . Within optimisation, Allgower and Georg [2] reviewed numerical continuation methods, of which homotopy is a special case. The contribution of the present paper is to use the deformation equation as an

153 *algorithmic* primitive—truncated, evaluated numerically, and applied uni-
 154 formly to identification and simulation—rather than as a device for analytic
 155 series construction.

156 3. The universal deformation equation

157 3.1. Liao’s zeroth-order deformation

158 Given a nonlinear operator $\mathcal{N}$ acting on an unknown ϕ with $\mathcal{N}[\phi^*] = 0$,
 159 the deformation equation (2) constructs a one-parameter family $\phi(q)$, $q \in$
 160 $[0, 1]$, such that $\phi(0) = \phi_0$ and $\phi(1) = \phi^*$. The freedom in choosing $\mathcal{L}$,
 161 ϕ_0 , and c_0 is the analytic content of the method [27]. We use this freedom
 162 algorithmically: $\mathcal{L}$ encodes prior knowledge, ϕ_0 is the best available cheap
 163 approximation, and c_0 scales the corrections.

164 3.2. The universal scalar corrections

165 Expanding $\phi(q) = \phi_0 + \sum_{m \geq 1} z_m q^m$ and Taylor-expanding $\mathcal{N}[\phi_0 + \delta]$
 166 about ϕ_0 , matching coefficients of q^m in (2) yields, in scalar form,

$$\boxed{z_1 = \gamma, \quad z_2 = \gamma(1 + \sigma), \quad z_3 = \gamma(1 + \sigma)^2 + \frac{c_0 \mathcal{N}_0''}{2\lambda} \gamma^2,} \quad (3)$$

167 where

$$\gamma \triangleq \frac{c_0 \mathcal{N}_0}{\lambda}, \quad \sigma \triangleq \frac{c_0 \mathcal{N}_0'}{\lambda}, \quad (4)$$

168 $\mathcal{N}_0 = \mathcal{N}[\phi_0]$, $\mathcal{N}_0' = \partial \mathcal{N} / \partial \phi|_{\phi_0}$, $\mathcal{N}_0'' = \partial^2 \mathcal{N} / \partial \phi^2|_{\phi_0}$, and λ is the scalar repre-
 169 sentation of $\mathcal{L}$. For finite-dimensional problems λ is a chosen positive scalar;
 170 for differential operators λ incorporates the discretisation. The truncated
 171 series converges if $|1 + \sigma| < 1$, equivalently $0 < \mathcal{N}_0' / \lambda < 2$ for $c_0 = -1$.

172 **Remark 1** (Newton and Halley as special cases). *With $\lambda = \mathcal{N}_0'$ and $c_0 = -1$
 173 one obtains $\sigma = -1$, hence $z_2 = 0$ and z_3 reduces to the classical Halley
 174 correction $-\frac{1}{2} \mathcal{N}_0^2 \mathcal{N}_0'' / (\mathcal{N}_0')^3$ [18]. The expressions (3) therefore generalise
 175 Newton–Halley to arbitrary λ , with the σ -parametrised one-parameter family
 176 interpolating between the two and beyond.*

177 **Remark 2** (Vectorial reading). *For $\phi \in \mathbb{R}^n$, the scalar $\mathcal{N}_0'$ is replaced by
 178 a Jacobian and $\mathcal{N}_0''$ by a Hessian tensor. With $\lambda = J^T J$ (Gauss–Newton
 179 metric) and a diagonal Hessian approximation, the expressions (3) produce
 180 the Levenberg–Marquardt step at order one and a Halley-type correction at
 181 order two. The diagonal Hessian approximation is a deliberate choice to
 182 retain $O(NM)$ cost per step; we discuss its accuracy in Section 4.*

183 4. Identification by homotopy

184 4.1. Problem statement and bilinear factorisation

185 We represent the unknown nonlinearity by a parametric approximator
186 with linear last layer,

$$\hat{f}(y) = \Phi(y; \boldsymbol{\eta}) \mathbf{w}, \quad (5)$$

187 where $\Phi : \mathbb{R} \rightarrow \mathbb{R}^{1 \times M}$ is a row of basis functions parametrised by $\boldsymbol{\eta}$ (centres
188 and widths for an RBF; hidden weights and biases for an MLP), and $\mathbf{w} \in$
189 $\mathbb{R}^M$ are linear weights. Given training data $\{(y_i, f_i)\}_{i=1}^N$, the identification
190 residual is

$$\mathcal{N}_{\text{id}}(\boldsymbol{\eta}, \mathbf{w}) = \Phi(\boldsymbol{\eta}) \mathbf{w} - \mathbf{f}, \quad (6)$$

191 where $\Phi_{ij} = \Phi(y_i; \boldsymbol{\eta})_j$ is the $N \times M$ design matrix and $\mathbf{f} = (f_1, \dots, f_N)^T$.

192 The residual is *linear* in $\mathbf{w}$ and *nonlinear* in $\boldsymbol{\eta}$. We exploit this by al-
193 ternating an exact linear solve on $\mathbf{w}$ with curvature-aware corrections on
194 $\boldsymbol{\eta}$.

195 4.2. Linear sub-problem (last-layer weights)

196 For fixed $\boldsymbol{\eta}$, the optimal weights are

$$\mathbf{w}^* = (\Phi^T \Phi)^{-1} \Phi^T \mathbf{f}. \quad (7)$$

197 This is the deformation correction z_1 with $\mathcal{L} = \Phi^T \Phi$ applied to the linear
198 residual; structurally, $\partial^2 \mathcal{N}_{\text{id}} / \partial \mathbf{w}^2 = 0$, so $z_2 = z_3 = 0$ and the solution is
199 exact in one step.

200 4.3. Nonlinear sub-problem (internal parameters)

201 After substituting $\mathbf{w} = \mathbf{w}^*(\boldsymbol{\eta})$, the residual is a function of $\boldsymbol{\eta}$ only. The
202 Jacobian $J_{\boldsymbol{\eta}}$ is computed analytically; for an RBF with centre c_j and width
203 λ_j ,

$$\frac{\partial \mathcal{N}_i}{\partial c_j} = w_j \varphi' \left(\frac{|y_i - c_j|}{\lambda_j} \right) \frac{\text{sgn}(y_i - c_j)}{\lambda_j}, \quad (8)$$

$$\frac{\partial \mathcal{N}_i}{\partial \lambda_j} = -w_j \varphi' \left(\frac{|y_i - c_j|}{\lambda_j} \right) \frac{|y_i - c_j|}{\lambda_j^2}. \quad (9)$$

204 For an MLP the corresponding derivatives are obtained by analytical dif-
205 ferentiation of the activation, e.g. $\sigma'(x) = \sigma(x)(1 - \sigma(x))$ and $\sigma''(x) =$
206 $\sigma'(x)(1 - 2\sigma(x))$ for the logistic sigmoid. We emphasise that these closed-form

207 derivatives *are* the parameter-by-parameter chain rule and so are mathemat-
 208 ically equivalent to backpropagation; the difference is downstream of the
 209 gradient. We do not use a learning rate or stochastic batches.

210 The Newton correction on the nonlinear parameters is

$$\Delta\boldsymbol{\eta}_1 = -(J_\eta^T J_\eta)^{-1} J_\eta^T \mathcal{N}_0, \quad (10)$$

211 which is z_1 of (3) read in vectorial form with $\lambda = J_\eta^T J_\eta$. For the Halley
 212 correction we use a *diagonal* Hessian approximation,

$$\Delta\eta_{j,2} \approx -\frac{1}{2} \frac{(\mathcal{N}_1^T \partial^2 \mathcal{N} / \partial \eta_j^2) (\Delta\eta_{j,1})^2}{(J_{\eta j}^T J_{\eta j})}, \quad (11)$$

213 where $\mathcal{N}_1$ is the residual after the Newton step.

214 *Diagonal-Hessian approximation: scope and effect.* Equation (11) retains
 215 only the j -diagonal entries of the parameter Hessian. The exact Halley step
 216 contains additional off-diagonal contributions that scale with the off-diagonal
 217 Hessian entries. The diagonal approximation is exact when the Hessian is
 218 diagonally dominant (in particular when the basis functions φ_j have non-
 219 overlapping supports relative to the data) and is otherwise a controlled ap-
 220 proximation that preserves quadratic convergence under the same conditions
 221 as Newton. We use it because (i) it keeps the per-iteration cost at $O(NM)$
 222 and (ii) it is the form that survives the universal scalar parametrisation (3).
 223 For problems in which the basis functions overlap strongly, a Kronecker-
 224 factored Hessian [32] can replace the diagonal in (11) without changing the
 225 rest of the framework.

226 4.4. Algorithm

227 Algorithm 1 states the identification procedure.

228 5. Simulation by integral homotopy

229 5.1. Operator splitting and analytic starting solution

230 After identification, the simulation problem is to integrate $\dot{y} + \hat{f}(y) = u(t)$
 231 from $y(0) = y_0$. We linearise the identified RHS at a reference point $\bar{y}$,

$$k = \hat{f}'(\bar{y}), \quad \tilde{f}(y) = \hat{f}(y) - k y, \quad (12)$$

Table 1: Identification algorithm (one outer iteration; iterate until $\|\mathcal{N}\| < \varepsilon$).

Step	Action
1	Initialise $\boldsymbol{\eta}^{(0)}$ (K-means for RBF; uniform grid for MLP).
2	Solve $\mathbf{w}^* \leftarrow (\boldsymbol{\Phi}^T \boldsymbol{\Phi})^{-1} \boldsymbol{\Phi}^T \mathbf{f}$ (z_1 exact, one step).
3	Form J_η analytically; apply Newton step (10).
4	Recompute $\mathbf{w}^*$ and the residual; apply Halley step (11).
5	Recompute $\mathbf{w}^*$; check $\ \mathcal{N}\ < \varepsilon$.

so that $\tilde{f}'(\bar{y}) = 0$ by construction, and split

$$\dot{y} + k y = u(t) - \tilde{f}(y). \quad (13)$$

The left-hand side is a first-order linear operator with analytic Green's function $K(t, s) = e^{-k(t-s)}$, yielding the starting solution

$$y_0(t) = y_0 e^{-kt} + \int_0^t K(t, s) u(s) ds. \quad (14)$$

5.2. Volterra integral residual

Equation (13) is equivalent to the Volterra integral equation

$$y(t) = y_0(t) - \int_0^t K(t, s) \tilde{f}(y(s)) ds. \quad (15)$$

Discretising the integral with the trapezoidal rule at step T and maintaining the convolution as an IIR accumulator S_n with update $S_n = e^{-kT} S_{n-1} + T w_n [u_n - \tilde{f}(y_n)]$ (where w_n are the trapezoidal weights), the per-step residual reads

$$g_{\text{sim}}(y_n) = y_n + T w_n \tilde{f}(y_n) - T w_n u_n - e^{-kT} S_{n-1} - y_0 e^{-kt_n}. \quad (16)$$

The Jacobian $g'_{\text{sim}} = 1 + T w_n \tilde{f}'(y_n)$ is $O(1)$ and contains no $1/T$ factor; sensor noise on u is multiplied by T rather than divided by T , which is the standard advantage of integral formulations [3, 24].

5.3. Per-step homotopy correction

The per-step deformation uses (3) with $\mathcal{N}_0 = g_{\text{sim}}(y_n^{(0)})$, $\mathcal{N}'_0 = g'_{\text{sim}}$, $\mathcal{N}''_0 = T w_n \tilde{f}''(y_n)$, and $\lambda = g'_{\text{sim}}$. With this choice $\sigma = -1$, $z_2 = 0$, and z_3 is the standard Halley correction. The cost is one linear-system solve (scalar, in the SISO case) and three evaluations of $\tilde{f}$ per step.

249 6. Unification: identification and simulation as one equation

Table 2: Subproblems in the unified framework. Each row is an instance of the deformation equation (2) with the indicated choice of $(\mathcal{L}, \mathcal{N}, \phi, \lambda)$.

Level	Sub-problem	$\mathcal{L}$	$\mathcal{N}$	ϕ	λ
1a	Last-layer weights	$\Phi^T \Phi$	$\Phi \mathbf{w} - \mathbf{f}$	$\mathbf{w} \in \mathbb{R}^M$	$\Phi^T \Phi$
1b	Internal parameters	$J_\eta^T J_\eta$	$\mathcal{N}_{\text{id}}(\boldsymbol{\eta})$	$\boldsymbol{\eta} \in \mathbb{R}^{2M}$	$J_\eta^T J_\eta$
2	ODE forward step	$\text{d/dt} + \frac{\cdot}{k}$	$\dot{\mathbf{y}} + \tilde{\mathbf{f}}(\mathbf{y}) - \mathbf{u}$	$\mathbf{y}_n \in \mathbb{R}$	$\mathbf{g}'_{\text{sim}}$

250 **Theorem 3** (Unified homotopy framework). *The three sub-problems in Ta-*
 251 *ble 2 are instances of the deformation equation (2), with corrections given by*
 252 *the universal scalar formulae (3). Specifically:*

- 253 (i) *Level 1a is solved exactly in one step ($z_2 = z_3 = 0$ structurally).*
- 254 (ii) *Level 1b converges if $0 < \mathcal{N}'_0/\lambda < 2$ at each iterate, with quadratic rate*
 255 *from z_1 alone and cubic rate from $z_1 + z_2$ when the Hessian is diagonally*
 256 *dominant.*
- 257 (iii) *Level 2 converges per step if $\tilde{\mathbf{f}}$ is Lipschitz and $T < 1/\max|\tilde{\mathbf{f}}'|$.*
- 258 (iv) *No scalar learning rate or stochastic batching is used at any level.*

259 *Proof sketch.* (i) Linearity of the residual in $\mathbf{w}$. (ii) Standard Newton con-
 260 vergence with curvature correction; the diagonal-dominance condition is suf-
 261 ficient for the diagonal Hessian approximation to retain Halley’s cubic rate.
 262 (iii) Compactness of the Volterra operator implies convergence of the Neu-
 263 mann series for finite horizons [24]; the step-size condition is the per-step
 264 Newton condition on the integral residual. (iv) By construction. $\square$

265 6.1. Recovery of classical methods

266 7. Constructive dimensioning

267 We now derive the dimensioning chain.

Table 3: The framework recovers classical numerical methods as parameter choices in (3).

Choice	Method recovered	Reference
$\mathcal{L} = \Phi^T \Phi$, z_1 on $\mathbf{w}$	Linear least squares	Björck [7]
$\mathcal{L} = J^T J$, z_1 on $(\mathbf{w}, \boldsymbol{\eta})$	Variable Projection / Gauss–Newton	Golub and Pereyra [17], Sjölund et al. [44]
$\mathcal{L} = J^T J$, $z_1 + z_2$ scalar	Halley’s method	Gutiérrez and Hernández [18]
$\mathcal{L} = \text{d/dt}$, integral residual	Picard iteration	Atkinson [3]
$\mathcal{L} = \text{d/dt} + k$, $z_1 + z_2 + z_3$ on integral residual	Integral Newton–Halley	This paper, Section 5

268 7.1. From simulation accuracy to identification accuracy

269 **Theorem 4** (Gronwall propagation). *Let $y(t)$ solve $\dot{y} + f(y) = u$ and $\hat{y}(t)$*
270 *solve $\dot{y} + \hat{f}(y) = u$ with the same initial condition. If f is Lipschitz with*
271 *constant L and $\|f - \hat{f}\|_\infty = \varepsilon_{\text{id}}$, then for $t \in [0, T_f]$,*

$$\|y - \hat{y}\|_{\infty, [0, T_f]} \leq \varepsilon_{\text{id}} \frac{e^{LT_f} - 1}{L}. \quad (17)$$

272 *Proof.* Standard Gronwall on $e(t) = y(t) - \hat{y}(t)$ with $|\dot{e}| \leq L|e| + \varepsilon_{\text{id}}$. □

273 Inverting,

$$\varepsilon_{\text{id}} \leq \varepsilon_{\text{sim}} \frac{L}{e^{LT_f} - 1}. \quad (18)$$

274 7.2. From identification accuracy to network size

275 **Proposition 5** (RBF approximation rate, 10). *Let $f \in C^2([a, b])$ with $\|f''\|_\infty =$*
276 *K_2 . For M Gaussian RBF centres with widths chosen by nearest-neighbour*
277 *scaling,*

$$\inf_{\mathbf{w}, \boldsymbol{\eta}} \|f - \hat{f}_M\|_\infty \leq C_\varphi \frac{K_2 (b - a)^2}{M^2}, \quad (19)$$

278 *with C_φ a basis-dependent constant ($C_\varphi = 1/8$ for equispaced centres).*

279 For sigmoid MLPs, the analogous rate is governed by the Barron norm
 280 $\|f\|_{\mathcal{B}}$ [4] with $M_{\text{MLP}} \propto \|f\|_{\mathcal{B}}^2 / \varepsilon_{\text{id}}^2$.

Corollary 6 (Minimum network size).

$$M_{\min} \geq (b - a) \sqrt{\frac{C_{\varphi} K_2 (e^{L T_f} - 1)}{L \varepsilon_{\text{sim}}}} \quad (\text{RBF}), \quad (20)$$

281 with the analogous Barron-norm expression for an MLP.

282 7.3. Data sufficiency and step-size

283 The Newton step (10) requires $J_{\eta} \in \mathbb{R}^{N \times 2M}$ to have full column rank
 284 with moderate conditioning. A sufficient condition is that the fill distance
 285 $h_X = \max_j \min_i |y_i - c_j|$ satisfy $h_X \lesssim \min_j \lambda_j$, which yields

$$N \geq \beta M, \quad \beta \geq 3. \quad (21)$$

286 The maximum admissible per-step size for the integral simulation is

$$T_{\max} < \frac{1}{\max_y |\tilde{f}'(y)|}, \quad (22)$$

287 which is minimised by choosing $k = \hat{f}'(\bar{y})$ at the operating average $\bar{y}$.

288 7.4. The dimensioning chain

289 **Theorem 7** (Constructive dimensioning). *Given ε_{sim} , T_f , L , K_2 , $[a, b]$:*

$$\boxed{\varepsilon_{\text{sim}} \xrightarrow{(18)} \varepsilon_{\text{id}} \xrightarrow{(20)} M_{\min} \xrightarrow{(21)} N_{\min} \xrightarrow{(22)} T_{\max}.} \quad (23)$$

290 *The chain is constructive: each link is computable from system quantities,*
 291 *and the homotopy algorithm of Section 4 attains the predicted ε_{id} when $M \geq$*
 292 *$M_{\min}$ and $N \geq N_{\min}$.*

293 8. The role of prior knowledge: operator-side inductive bias

294 The dimensioning bound (20) depends on $K_2 = \|\tilde{f}''\|_{\infty}$, the curvature
 295 of the *residual* nonlinearity left after the linear operator $\mathcal{L}$ absorbs what is
 296 already known. The choice of $\mathcal{L}$ therefore acts as a quantitative inductive
 297 bias.

Table 4: Operator choice, residual nonlinearity, and required network size. The last row is the limit of full prior knowledge.

Prior knowledge	$\mathcal{L}$	$\tilde{f} = f - \mathcal{L}^{-1}[f]$	$M_{\min}$
None	I (identity)	f (entire)	Largest
Linear approximation	$d/dt + k$	$f(y) - k y$	Smaller
Quadratic approximation	$d/dt + k + k_2 y$	$f(y) - k y - k_2 y^2$	Smaller still
Exact model	$\mathcal{L} =$ closed-form f	$\tilde{f} = 0$	0

298 *PINN contrast.* PINNs [38] place physics in the loss as a soft penalty $\mu \|\dot{\hat{y}} +$
 299 $f(\hat{y}) - u\|^2$. The network architecture is unchanged by the addition; the
 300 physics term constrains the optimisation landscape but does not reduce the
 301 function class the network must span. In our framework physics enters $\mathcal{L}$ and
 302 therefore $\tilde{f}$. The required M shrinks as $\|\tilde{f}''\|_\infty$ shrinks. The two approaches
 303 are complementary rather than alternative: a PINN that uses a homotopy-
 304 style operator split for its residual can in principle inherit both benefits.

305 *Architectures with bilinear last layers.* Table 5 summarises common neural
 306 architectures from the standpoint of the bilinear factorisation (5). We mark
 307 explicitly which architectures are validated experimentally in this paper and
 308 which are listed only as theoretical extensions consistent with the bilinear
 309 structure; the latter are not claimed as supported results.

310 9. Numerical demonstrations

311 9.1. Benchmark

312 We use the orifice-tank discharge model

$$\dot{h} + 0.5\sqrt{h} = u(t), \quad h(0) = 1.0, \quad (24)$$

313 on $t \in [0, 10]$ with a step-ramp excitation: $u(t) = 0$ for $t < 2$; $u(t) =$
 314 $0.075(t - 2)$ for $t \in [2, 6]$; $u(t) = 0.3$ for $t \geq 6$. Twenty data points are
 315 generated by RK45 with relative tolerance 10^{-10} . The unknown nonlinearity
 316 $f(h) = 0.5\sqrt{h}$ is identified and the ODE is then forward-simulated with the
 317 identified model.

Table 5: Common neural architectures expressed as $\hat{f} = \Phi(y; \eta) \mathbf{w}$ with linear last-layer weights and nonlinear internal parameters. The third column indicates the empirical status *in this paper*; theoretical extensions are not validated here.

Arch.	$\hat{f}$	Linear params $\mathbf{w}$	Status here
RBF	$\sum_j w_j \varphi(\ y - c_j\ /\lambda_j)$	$\{w_j\}$	Demonstrated
MLP (1 hidden)	$W_2 \sigma(W_1 y + b_1) + b_2$	$\{W_2, b_2\}$	Demonstrated
CNN (last layer)	$W_{\text{fc}} \cdot \text{conv}(y; \mathcal{K})$	W_{fc}	Theoretical extension
LSTM	$W_o \cdot h_T(y; W_h)$	W_o	Theoretical extension
Transforme head	$W_{\text{head}} \cdot \text{attn}(y; Q, K, V)$	W_{head}	Theoretical extension

318 9.2. Architecture A: RBF with five Gaussian centres

319 *Identification.* K-means initialises the centres $\mathbf{c}^{(0)}$; nearest-neighbour dis-
320 tances initialise the widths $\lambda^{(0)}$. One outer iteration of Algorithm 1 reduces
321 the residual from $\|\mathcal{N}_0\| = 0.112$ to $\|\mathcal{N}_f\| = 3.1 \cdot 10^{-2}$. Figure 1 shows the
322 identified function and the basis components.

323 *Simulation.* With $k = \hat{f}'(\bar{h}) = 0.654$ and 200 trapezoidal steps, the integral
324 homotopy yields a maximum simulation error of 2.3% over the full hori-
325 zon, against 82% for a linearised baseline (factor 35 $\times$). Figure 2 shows the
326 forward-simulated trajectory.

327 9.3. Architecture B: MLP with eight sigmoid units

328 *Identification.* For this benchmark the step-ramp occupies $t \in [1, 3]$ ($u(t) =$
329 $0.15(t - 1)$, held at 0.3 thereafter), producing a 20-point trajectory analo-
330 gous to that of Architecture A. Sigmoid centres are uniformly distributed on
331 $[h_{\min}, h_{\max}]$ with initial slope $W_1^{(0)} = 5$ plus small seeded perturbations. At
332 this initialisation the exact linear-stage solve (Step 2 of Algorithm 1) already
333 drives the residual to machine precision $\|\mathcal{N}_f\| = 1.7 \cdot 10^{-8}$, so the outer loop
334 terminates after a single iteration without invoking the nonlinear corrections.




Figure 1: RBF identification on the orifice-tank data: identified function (solid), true $f(h) = 0.5\sqrt{h}$ (dashed), and 20 training points. Five Gaussian centres adapted by K-means + homotopy.

335 *Simulation.* A second-order explicit integrator combined with the same per-
 336 step homotopy correction yields a maximum simulation error of 0.96%, against
 337 7.7% for a linearised baseline (factor 8.1 $\times$). Figure 3 compares the two ar-
 338 chitectures on the same benchmark.

339 9.4. Comparative summary

340 9.5. Extension to coupled 2D systems: Lotka-Volterra predator-prey

341 To demonstrate that the framework generalises to multi-dimensional cou-
 342 pled systems, we apply it to the Lotka-Volterra predator-prey model,

$$\dot{x} = x(a - by), \quad \dot{y} = y(-c + dx), \quad (25)$$

343 with parameters $(a, b, c, d) = (1.0, 0.5, 1.0, 0.5)$ and initial conditions $(x_0, y_0) =$
 344 $(1.5, 1.0)$. The system exhibits oscillatory dynamics with a period of approx-
 345 imately $T \approx 6.5$ time units. We generate 40 training samples over $t \in [0, 15]$
 346 (roughly 2.3 periods) and identify the coupled nonlinearities $f_x(x, y) = x(a -$
 347 $by)$ and $f_y(x, y) = y(-c + dx)$ using two independent 2D Gaussian RBF ap-
 348 proximators, each with $M = 16$ centres.




Figure 2: RBF forward simulation of the orifice-tank ODE with the identified model: reference trajectory (solid), homotopy-simulated trajectory (dashed), and linearised baseline (dotted). Maximum error 2.3% vs. 82%.

349 *Identification.* The identification is performed in two decoupled stages: (i) fit
 350 f_x from the samples $\{(x_i, y_i, \dot{x}_i)\}$, where $\dot{x}_i$ is estimated by finite differences;
 351 (ii) fit f_y from $\{(x_i, y_i, \dot{y}_i)\}$. Each RBF uses a linear least-squares solve on
 352 the weights (exact one-step solution), with centres initialised on a uniform
 353 2D grid over the observed (x, y) domain. The identification residuals are
 354 $\|\mathcal{N}_x\| = 3.6 \cdot 10^{-1}$ and $\|\mathcal{N}_y\| = 2.4 \cdot 10^{-1}$, reflecting the higher complexity of
 355 the coupled 2D nonlinearities compared to the 1D orifice-tank case.

356 *Forward simulation.* Figure 4 shows the identified trajectories over a 15-time-
 357 unit horizon (400 integration steps with RK45). The phase portrait (panel a)
 358 demonstrates that the identified model preserves the qualitative oscillatory
 359 structure of the predator-prey cycle, though with quantitative drift accumu-
 360 lating over multiple periods. Time series (panels b and c) show maximum
 361 errors of 1.73 in prey population and 2.75 in predator population (absolute
 362 values), corresponding to relative errors of approximately 25% and 40% at
 363 the peaks, respectively. These errors are larger than in the 1D case due to:

Comparison: MLP vs Linear vs RBF
(* RBF values from demo_21paper.py)

Metric	MLP (8 sigmoids)	Linear	RBF (5 centers)
Total parameters	25	1	15
Nonlinear params	16 (W1, b1)	0	10 (c, λ)
Linear params	9 (W2, b2)	1 (w)	5 (w)
Outer iterations	1	1	1
Max sim error	0.0096	0.0771	0.023*
Error reduction	8.1×	-	35×

Figure 3: Architecture comparison on the orifice-tank benchmark: 5-centre RBF and 8-unit sigmoid MLP, each identified from 20 data points of its orifice-tank trajectory with the same homotopy algorithm. The MLP attains a smaller identification residual and simulation error at the cost of additional parameters.

364 (i) the coupled nature of the dynamics, where errors in one variable propagate
365 to the other; (ii) the oscillatory regime, where small phase shifts accumulate
366 over periods; (iii) the use of simple uniform-grid initialisation, chosen for re-
367 producibility, rather than adaptive clustering of the centres. Despite these
368 limitations, the identified model captures the essential predator-prey cycle
369 and demonstrates that the bilinear homotopy framework extends to multi-
370 dimensional coupled systems.

371 *Remark.* The Lotka-Volterra demonstration establishes the *proof of extensi-*
372 *bility*: the bilinear identification primitive (exact linear-stage solve, with the
373 closed-form Newton + Halley corrections available on the nonlinear param-
374 eters) extends to coupled multi-dimensional systems with no modifications
375 beyond fitting multiple independent approximators for each component func-

Table 6: RBF and MLP on the orifice-flow benchmark (each with its step-ramp excitation as described above), same homotopy algorithm.

Metric	RBF (5 centres)	MLP (8 sigmoids)
Total parameters	15	25
Linear parameters	5	9
Nonlinear parameters	10	16
Identification residual	$3.1 \cdot 10^{-2}$	$1.7 \cdot 10^{-8}$
Maximum simulation error	2.3%	0.96%
Improvement over linear baseline	$35\times$	$8.1\times$
Outer iterations	1	1
Learning rate	none	none

tion. In this 2D demonstration only the linear stage is exercised. Tighter accuracy would require: (i) adaptive centre placement via K-means or the Newton–Halley corrections on the nonlinear RBF parameters; (ii) higher-order correctors (z_3 , z_4) or iterative refinement; (iii) tailored integration schemes that exploit the Hamiltonian structure of conservative systems like Lotka-Volterra. Such refinements fall outside the scope of the present proof-of-concept regime but are natural follow-up directions.

9.6. Wall-clock comparison

To place the homotopy corrections in numerical context we run four gradient-based alternatives on the same MLP identification problem of Section 9 (Architecture B) with the same 20 training points and identical initialisation, on a single CPU core: (i) vanilla gradient descent over 5,000 iterations; (ii) Adam [1] ($\beta_1 = 0.9$, $\beta_2 = 0.999$) with a cosine learning-rate decay over 5,000 iterations; (iii) L-BFGS-B quasi-Newton method [12]; (iv) Trust-Region Reflective (TRF) least-squares [6], run to default tolerances with maximum 2,500 function evaluations. To avoid biasing the comparison through a hand-picked step size, the learning rates of the two first-order baselines are the best found by a grid sweep over $\eta \in \{10^{-4}, \dots, 10^{-1}\}$ under the identical 5,000-iteration budget (both select $\eta = 10^{-1}$; the sweep script `tune_optimizer_baselines.py` is included in the public repository). Median wall-clock times over five repetitions are reported in Table 7; the bench-




Figure 4: Lotka-Volterra predator-prey system identified via 2D RBF (exact linear-stage solve of the bilinear homotopy framework). (a) Phase portrait showing the oscillatory cycle (true trajectory: solid black; identified: dashed red; training data: blue circles). (b) Prey population $x(t)$ over time. (c) Predator population $y(t)$ over time. The identified model preserves the qualitative oscillatory structure over ~ 2.3 periods, with maximum absolute errors of 1.73 (prey) and 2.75 (predator).

397 mark script is included in the public repository (`demo_optimizer_comparison.py`).

Table 7: Wall-clock comparison on the MLP identification problem (8 sigmoids, 20 data points, identical initialisation, single CPU core, no warm start). Median over 5 runs.

Method	Final residual	Iter./nfev	Wall time
Gradient descent (tuned $\eta = 10^{-1}$)	$1.1 \cdot 10^{-2}$	5,000	0.25 s
Adam (cosine decay, tuned $\eta = 10^{-1}$)	$1.3 \cdot 10^{-3}$	5,000	0.31 s
L-BFGS-B (SciPy)	$2.5 \cdot 10^{-3}$	17 / 20	0.0025 s
Trust-Region Reflective	$3.1 \cdot 10^{-6}$	2,500	1.50 s
Homotopy ($z_1 + z_2$, this work)	$1.7 \cdot 10^{-8}$	1 outer	0.0002 s

398 Even with a tuned learning rate, vanilla gradient descent stalls at a resid-
399 ual of order 10^{-2} after 5,000 iterations, and tuned Adam with cosine decay
400 reaches $1.3 \cdot 10^{-3}$. L-BFGS-B exploits second-order information and con-
401 verges rapidly (20 function evaluations, 0.0025 s) to a comparable residual
402 of $2.5 \cdot 10^{-3}$. TRF achieves the smallest residual among the gradient-based
403 methods ($3.1 \cdot 10^{-6}$) but exhausts its function-evaluation budget (2,500) in
404 1.50 s. The homotopy method, by solving the linear sub-problem exactly and
405 applying curvature-aware corrections only to the lower-dimensional nonlinear

parameters, reaches $1.7 \cdot 10^{-8}$ (the Section 9 result) in one order of magnitude less wall-clock time than L-BFGS-B and four orders less than TRF. The advantage is structural: the bilinear separation eliminates the need for iterative convergence on the linear weights, leaving only the $2M$ nonlinear parameters (slopes and biases) for the Newton + Halley corrections.

9.7. Robustness to measurement noise

To assess practical applicability beyond idealized synthetic data, we repeat the orifice-tank identification under additive Gaussian noise. Training samples $\{h_i\}$ are corrupted by $h_i \leftarrow h_i + \epsilon_i$, where $\epsilon_i \sim \mathcal{N}(0, \sigma^2)$ and σ is chosen to yield relative noise levels of 1%, 5%, and 10% (measured against the standard deviation of the clean signal). For each noise level, we average results over five independent trials.

Table 8 reports identification residuals and forward-simulation errors for the 5-centre RBF under increasing noise, comparing the homotopy method with L-BFGS-B [12] (same initialization, all $3M$ parameters optimised jointly, maximum 5,000 iterations).

Table 8: Robustness to measurement noise: 5-centre RBF on orifice-tank, averaged over 5 trials per noise level.

Noise Level	Homotopy		L-BFGS-B	
	ID residual	Sim error	ID residual	Sim error
0% (clean)	$1.6 \cdot 10^{-2}$	0.010	$1.2 \cdot 10^{-2}$	0.009
1%	$2.0 \cdot 10^{-2}$	0.011	$1.7 \cdot 10^{-2}$	0.010
5%	$7.1 \cdot 10^{-2}$	0.018	$6.9 \cdot 10^{-2}$	0.017
10%	$1.4 \cdot 10^{-1}$	0.034	$1.4 \cdot 10^{-1}$	0.032

Figure 5 plots identification and simulation errors as functions of noise level. Both methods degrade gracefully: errors scale approximately linearly with noise in the moderate regime (1%–5%), and forward-simulation accuracy remains below 3.5% maximum error even at 10% noise. The homotopy method maintains performance comparable to L-BFGS-B across all noise levels, confirming that the Newton + Halley corrections are not overly sensitive to measurement uncertainty when applied to the bilinear structure.




Figure 5: Robustness to measurement noise on the orifice-tank identification and simulation. (a) Identification residual (L2 norm) vs noise level. (b) Maximum forward-simulation error over 200-step horizon vs noise level. Both homotopy and L-BFGS-B degrade gracefully, with comparable performance across all noise levels.

9.8. Verification of the dimensioning chain

For the orifice problem, $K_2 = \max |f''| = 0.5/(4h_{\min}^{3/2}) \approx 4$ on $[0.1, 1.5]$, $L \approx 1.6$, $T_f = 10$, target $\varepsilon_{\text{sim}} = 0.05$. The bound (20) with equispaced centres ($C_\varphi = 1/8$) gives $M_{\min} \geq 29$. The demonstration achieves $\varepsilon_{\text{sim}} \approx 0.023$ with $M = 5$ adaptive (K-means + homotopy) centres.

Interpretation of the discrepancy. The factor $\sim 5.8\times$ between the theoretical minimum $M_{\min} = 29$ and the empirically sufficient $M = 5$ reflects the conservative nature of worst-case approximation bounds. The bound (20) assumes:

- (i) Uniform approximation over the full domain $[h_{\min}, h_{\max}]$ with no exploitation of the smoothness or specific structure of $f(h) = 0.5\sqrt{h}$.
- (ii) Equispaced centre placement with fixed separation constant $C_\varphi = 1/8$, which is a conservative lower bound for general RBF configurations.
- (iii) Worst-case training data distribution, with no assumption that samples are concentrated in regions of higher curvature or nonlinearity.

In practice, K-means clustering adapts centre placement to the empirical data distribution, the homotopy corrections refine the nonlinear parameters beyond the initialisation, and the target function $0.5\sqrt{h}$ is smoother than a generic C^2 function with the same K_2 bound. Each of these factors contributes to the gap between the theoretical sufficient condition and the observed performance.

450 *Role of the dimensioning chain.* The chain should therefore be interpreted
 451 as a *feasibility certificate*—guaranteeing that a network of at least the pre-
 452 scribed size *exists* and will achieve the target accuracy under the stated
 453 assumptions—rather than a tight design rule or an exact predictor of the
 454 minimum required dimension. Tighter, problem-specific bounds would re-
 455 quire refined constants (e.g., Sobolev or Hölder norms instead of point-wise
 456 K_2 , data-dependent placement constants) or adaptive mesh estimates that
 457 depend on the identified function itself. Such refinements fall outside the
 458 scope of the present universal framework, which prioritises transparency and
 459 computability of the bound over tightness.

460 *When are the bounds informative?* The bounds are most informative when:

- 461 (i) The target function f has no exploitable structure beyond the assumed
 462 regularity (C^k smoothness).
- 463 (ii) Centre placement is constrained to be non-adaptive (e.g., uniform grids
 464 for real-time implementation).
- 465 (iii) A certificate of feasibility is required before training (e.g., embedded
 466 systems with fixed memory constraints).

467 Conversely, when adaptive placement, iterative refinement, or problem-specific
 468 structure can be exploited, the bounds provide an upper estimate of the
 469 required network size, and smaller architectures may suffice. The present
 470 orifice-tank demonstration exemplifies the latter regime: the bound is con-
 471 servative by design, and adaptive methods achieve the target with fewer
 472 resources.

473 10. Limitations and scope

474 The empirical scope of this paper is deliberately narrow and we state it
 475 explicitly here.

476 *Single benchmark.* The numerical study uses one one-dimensional first-order
 477 ODE. Multi-dimensional, stiff, and chaotic systems are not covered. The
 478 dimensioning chain (Theorem 7) extends formally to multi-dimensional states
 479 by replacing K_2 with a bound on the relevant tensor norm of $\tilde{f}$, but we have
 480 not validated the constants empirically.

481 *Two architectures.* Only RBF and MLP are demonstrated. The bilinear
 482 factorisation (5) applies to convolutional and recurrent architectures with a
 483 linear last layer, but the constants in the approximation rate (Proposition 5)
 484 and the conditioning of J_η depend on the architecture and are not analysed
 485 for these cases. The corresponding rows of Table 5 are explicitly marked as
 486 theoretical extensions, not as validated results.

487 *Diagonal Hessian approximation.* The Halley correction in (11) uses a diago-
 488 nal Hessian. This is exact for diagonally dominant Hessians and a controlled
 489 approximation otherwise. For basis functions with strongly overlapping sup-
 490 ports the approximation weakens. Replacing the diagonal by a Kronecker-
 491 factored Hessian [32] is a direct extension that we have not pursued here.

492 *Online refinement.* An online-refinement variant of the framework (recursive
 493 update of $\boldsymbol{\eta}$ as new data arrives) follows formally from the same deformation
 494 equation applied to a sliding tracking residual. We omit it from the present
 495 paper because we do not have empirical results to support it; it is the natural
 496 target of a follow-up study.

497 *Wall-clock comparison.* The wall-clock results in Table 7 cover one problem
 498 on one machine. The four baselines (tuned gradient descent and Adam on
 499 the full parameter vector, L-BFGS-B, and SciPy’s TRF) do not exhaust the
 500 state of the art on this problem class: K-FAC and trust-region Newton with
 501 analytic Hessian would close part of the gap. The point of the comparison is
 502 to show that exploiting the bilinear separation *at all*, rather than treating the
 503 network as a generic nonlinear least-squares object, eliminates the need for a
 504 learning rate or trust-region tuning on the small-data regime of interest here.
 505 A fairer comparison against curvature-aware baselines on larger problems is
 506 left for future work.

507 11. Relationship to companion work

508 The present manuscript is part of a programme of work by the authors
 509 on homotopy-based methods for identification and control. We clarify the
 510 delimitation explicitly here.

511 *Integral homotopy regressors.* A separate manuscript by the authors with
 512 S. J. Liao [39] develops the integral formulation of the discrete homotopy
 513 regressor as a Volterra operator with arbitrary compact kernel; the present

514 paper specialises that framework to the exponential kernel $K(t, s) = e^{-k(t-s)}$
 515 that arises naturally from operator-splitting at a linearisation point, and
 516 applies it inside the simulation level.

517 *Structural inductive bias in implicit models.* A separate manuscript by the
 518 authors [40] analyses the sample-complexity gain of *implicit* formulations
 519 $F(y, \hat{f}(y)) = 0$ via Rademacher complexity. The present paper analyses the
 520 network-size gain of *operator-split* formulations $\dot{y} + ky + \tilde{f}(y) = u$ via ap-
 521 proximation theory and the dimensioning chain. The two analyses address
 522 different mechanisms (generalisation versus approximation) and are comple-
 523 mentary; neither subsumes the other.

524 *Coupled-system identification.* A further manuscript [41] extends homotopy
 525 identification to coupled nonlinear ODEs with embedded MLP coupling func-
 526 tions. The bilinear factorisation and the linear-subproblem-exact-solve are
 527 shared with the present paper; the coupled case requires additional treatment
 528 of the nonlinear coupling that is specific to that setting.

529 12. Discussion

530 *Why closed-form corrections rather than gradient descent?* For the small
 531 networks and bilinear residuals we consider, the Jacobian and the diagonal
 532 of the parameter Hessian are cheap to compute and analytically available.
 533 Newton with curvature correction (Halley) converges quadratically and cu-
 534 bically respectively, and contains no scalar parameter that requires tuning.
 535 Gradient descent retains advantages for very large networks where the cost of
 536 forming $J^T J$ becomes prohibitive; the present framework targets the regime
 537 where it does not.

538 *Why the integral residual?* The differential residual $g_{\text{BDF}} \sim (y_n - y_{n-1})/T +$
 539 $\hat{f}(y_n) - u_n$ contains $1/T$, which amplifies sensor noise on u by the same factor.
 540 The integral residual (16) contains T multiplicatively and attenuates the
 541 same noise. Both have $O(T^2)$ discretisation error, but the noise propagation
 542 differs structurally.

543 *Where does prior knowledge help?* The operator-split form (13) encodes
 544 prior knowledge in k , the linearisation slope. Better prior knowledge (smaller
 545 residual nonlinearity) reduces $\|\tilde{f}''\|_\infty$ and therefore $M_{\min}$. This is the operator-
 546 side counterpart of the loss-side mechanism used by PINNs.

547 13. Conclusions

548 We have shown that Liao’s homotopy deformation equation provides a
 549 single algorithmic primitive that solves both nonlinear system identification
 550 and ODE simulation when the function approximator has a bilinear last-
 551 layer-linear / internal-nonlinear structure. The closed-form corrections z_1
 552 (Newton) and z_2 (Halley) read off the deformation equation play three coor-
 553 dinated roles: they solve the linear last-layer sub-problem exactly in one step,
 554 they correct the nonlinear internal parameters with curvature information,
 555 and they advance the integral ODE residual one step at a time. No scalar
 556 learning rate is required.

557 The principal theoretical contribution is the constructive dimensioning
 558 chain $\varepsilon_{\text{sim}} \rightarrow \varepsilon_{\text{id}} \rightarrow M_{\text{min}} \rightarrow N_{\text{min}} \rightarrow T_{\text{max}}$, which links a target simulation
 559 accuracy to network size, training-data size, and integration step. The chain
 560 admits an interpretation in which the choice of the auxiliary linear operator
 561 acts as a quantitative inductive bias: prior physical knowledge encoded in $\mathcal{L}$
 562 shrinks the residual nonlinearity that the network has to approximate and
 563 therefore the required network size.

564 The empirical study covers a one-dimensional first-order benchmark with
 565 two architectures (RBF, MLP). The reported maximum simulation errors
 566 (2.3% and 0.96%) and identification residuals (down to machine precision in
 567 one outer iteration) are consistent with the dimensioning predictions. Gen-
 568 eralisation to multi-dimensional, stiff, and chaotic dynamics, and to convo-
 569 lutional, recurrent, and attention-based architectures, is consistent with the
 570 bilinear structure but requires further work, which we have outlined explicitly
 571 in Section 10.

572 Code and data availability

573 Reference implementations of both demonstrations (RBF and MLP), to-
 574 gether with smoke tests, pinned dependencies, and figure-generation scripts,
 575 are publicly available:

- 576 • **GitHub:** <https://github.com/rodrigo1964-ai/unified-homotopy-framework>
- 577 • **Figshare archive:** <https://doi.org/10.6084/m9.figshare.31955865>
- 578 • **ORCID (R. H. Rodrigo):** <https://orcid.org/0000-0002-8787-0038>

579 Acknowledgements

580 This work was carried out at INAUT (UNSJ-CONICET). The authors
581 thank Shijun Liao for foundational work on the Homotopy Analysis Method
582 on which this paper builds.

583 References

- 584 [1] Kingma, D.P., Ba, J., 2015. Adam: A method for stochastic optimiza-
585 tion. In: *Proc. 3rd International Conference on Learning Representa-*
586 *tions (ICLR)*.
- 587 [2] Allgower, E.L., Georg, K., 2003. *Introduction to Numerical Continua-*
588 *tion Methods*. SIAM.
- 589 [3] Atkinson, K.E., 1997. *The Numerical Solution of Integral Equations of*
590 *the Second Kind*. Cambridge University Press.
- 591 [4] Barron, A.R., 1993. Universal approximation bounds for superpositions
592 of a sigmoidal function. *IEEE Trans. Inform. Theory* 39, 930–945.
- 593 [5] Becker, S., LeCun, Y., 1988. Improving the convergence of back-
594 propagation learning with second order methods. In: *Proc. 1988 Con-*
595 *nectionist Models Summer School*, pp. 29–37.
- 596 [6] Branch, M.A., Coleman, T.F., Li, Y., 1999. A subspace, interior, and
597 conjugate gradient method for large-scale bound-constrained minimiza-
598 tion problems. *SIAM J. Sci. Comput.* 21, 1–23.
- 599 [7] Björck, Å., 1996. *Numerical Methods for Least Squares Problems*. SIAM.
- 600 [8] Broomhead, D.S., Lowe, D., 1988. Multivariable functional interpola-
601 tion and adaptive networks. *Complex Systems* 2, 321–355.
- 602 [9] Brunton, S.L., Proctor, J.L., Kutz, J.N., 2016. Discovering governing
603 equations from data by sparse identification of nonlinear dynamical sys-
604 tems. *Proc. Natl. Acad. Sci. USA* 113, 3932–3937.
- 605 [10] Buhmann, M.D., 2003. *Radial Basis Functions: Theory and Implemen-*
606 *tations*. Cambridge University Press.

- 607 [11] Butcher, J. C., 2016. *Numerical Methods for Ordinary Differential Equations*, 3rd ed. Wiley.
- 608
- 609 [12] Byrd, R. H., Lu, P., Nocedal, J., Zhu, C., 1995. A limited memory algorithm for bound constrained optimization. *SIAM J. Sci. Comput.* 16,
- 610 1190–1208.
- 611
- 612 [13] Champion, K., Lusch, B., Kutz, J. N., Brunton, S. L., 2019. Data-driven discovery of coordinates and governing equations. *Proc. Natl. Acad. Sci. USA* 116, 22445–22451.
- 613
- 614
- 615 [14] Chen, S., Cowan, C. F. N., Grant, P. M., 1991. Orthogonal least squares learning algorithm for radial basis function networks. *IEEE Trans. Neural Networks* 2, 302–309.
- 616
- 617
- 618 [15] Chen, R. T. Q., Rubanova, Y., Bettencourt, J., Duvenaud, D., 2018. Neural ordinary differential equations. In: *Advances in Neural Information Processing Systems (NeurIPS)*.
- 619
- 620
- 621 [16] Cybenko, G., 1989. Approximation by superpositions of a sigmoidal function. *Math. Control Signals Syst.* 2, 303–314.
- 622
- 623 [17] Golub, G. H., Pereyra, V., 1973. The differentiation of pseudo-inverses and nonlinear least squares problems whose variables separate. *SIAM J. Numer. Anal.* 10, 413–432.
- 624
- 625
- 626 [18] Gutiérrez, J. M., Hernández, M. A., 1997. A family of Chebyshev–Halley type methods in Banach spaces. *Bull. Aust. Math. Soc.* 55, 113–130.
- 627
- 628 [19] Hairer, E., Nørsett, S. P., Wanner, G., 1993. *Solving Ordinary Differential Equations I: Nonstiff Problems*, 2nd ed. Springer.
- 629
- 630 [20] Haykin, S., 2009. *Neural Networks and Learning Machines*, 3rd ed. Prentice Hall.
- 631
- 632 [21] Hornik, K., 1991. Approximation capabilities of multilayer feedforward networks. *Neural Networks* 4, 251–257.
- 633
- 634 [22] Kaheman, K., Kutz, J. N., Brunton, S. L., 2020. SINDy-PI: a robust algorithm for parallel implicit sparse identification of nonlinear dynamics. *Proc. R. Soc. A* 476, 20200279.
- 635
- 636

- [23] Karniadakis, G.E., Kevrekidis, I.G., Lu, L., Perdikaris, P., Wang, S., Yang, L., 2021. Physics-informed machine learning. *Nat. Rev. Phys.* 3, 422–440.
- [24] Kress, R., 2014. *Linear Integral Equations*, 3rd ed. Springer.
- [25] Krishnapriyan, A.S., Gholami, A., Zhe, S., Kirby, R., Mahoney, M., 2021. Characterizing possible failure modes in physics-informed neural networks. In: *Advances in Neural Information Processing Systems*.
- [26] Liao, S.J., 2003. *Beyond Perturbation: Introduction to the Homotopy Analysis Method*. Chapman & Hall/CRC.
- [27] Liao, S.J., 2010. An optimal homotopy-analysis approach for strongly nonlinear differential equations. *Commun. Nonlinear Sci. Numer. Simul.* 15, 2003–2016.
- [28] Liao, S.J., 2012. *Homotopy Analysis Method in Nonlinear Differential Equations*. Springer.
- [29] Lukoševičius, M., Jaeger, H., 2009. Reservoir computing approaches to recurrent neural network training. *Computer Science Review* 3, 127–149.
- [30] Ljung, L., 1999. *System Identification: Theory for the User*, 2nd ed. Prentice Hall.
- [31] Martens, J., 2010. Deep learning via Hessian-free optimization. In: *Proc. 27th Int. Conf. Machine Learning (ICML)*.
- [32] Martens, J., Grosse, R., 2015. Optimizing neural networks with Kronecker-factored approximate curvature. In: *Proc. 32nd Int. Conf. Machine Learning (ICML)*.
- [33] Narendra, K.S., Parthasarathy, K., 1990. Identification and control of dynamical systems using neural networks. *IEEE Trans. Neural Networks* 1, 4–27.
- [34] Nelles, O., 2001. *Nonlinear System Identification*. Springer.
- [35] Park, J., Sandberg, I.W., 1991. Universal approximation using radial-basis-function networks. *Neural Computation* 3, 246–257.

- [36] Poggio, T., Girosi, F., 1990. Networks for approximation and learning. *Proc. IEEE* 78, 1481–1497.
- [37] Rackauckas, C., Ma, Y., Martensen, J., Warner, C., Zubov, K., Supekar, R., Skinner, D., Ramadhan, A., Edelman, A., 2020. Universal differential equations for scientific machine learning. arXiv:2001.04385.
- [38] Raissi, M., Perdikaris, P., Karniadakis, G.E., 2019. Physics-informed neural networks: A deep learning framework for solving forward and inverse problems involving nonlinear partial differential equations. *J. Comput. Phys.* 378, 686–707.
- [39] Rodrigo, R.H., Patiño, H.D., Liao, S.J., 2026a. Integral homotopy regressors: unifying continuous and discrete HAM via compact Volterra operators. Manuscript in preparation.
- [40] Rodrigo, R.H., Patiño, H.D., 2026b. Structural inductive bias reduces sample complexity in implicit neural models for dynamical systems. Manuscript in preparation.
- [41] Rodrigo, R.H., Patiño, H.D., 2026c. Homotopy identification of coupled nonlinear systems with embedded MLP coupling functions. Manuscript in preparation.
- [42] Rumelhart, D.E., Hinton, G.E., Williams, R.J., 1986. Learning representations by back-propagating errors. *Nature* 323, 533–536.
- [43] Sjöberg, J., Zhang, Q., Ljung, L., et al., 1995. Nonlinear black-box modeling in system identification: a unified overview. *Automatica* 31, 1691–1724.
- [44] Sjölund, J., Malleshham, B., LaConte, S.M., Niazi, K., Madsen, K.H., Sander, O., Andersen, L.N., 2014. Variable projection methods in neural-network parameter estimation. *Numer. Linear Algebra Appl.* 21, 657–684.
- [45] Turkyilmazoglu, M., 2010. A note on the homotopy analysis method. *Appl. Math. Lett.* 23, 1226–1230.
- [46] Vajravelu, K., Van Gorder, R.A., 2013. *Nonlinear Flow Phenomena and Homotopy Analysis*. Springer.

- 697 [47] Wang, S., Teng, Y., Perdikaris, P., 2021. Understanding and mitigating
698 gradient pathologies in physics-informed neural networks. *SIAM J. Sci.*
699 *Comput.* 43, A3055–A3081.